

SEMANTIC DATA CATALOGUE

Introduction

The Semantic Catalogue aims to unify search across the [Consumer Data Research Centre \(CDRC\)](#), [Administrative Data Research UK \(ADRUk\)](#), [UK Data Service \(UKDS\)](#), and [Natural Environment Research Council \(NERC\)](#) data catalogues. To improve the discoverability of data held within these catalogues, this system implements *semantic search*, moving beyond traditional search methods, which focus solely on the presence of keywords. Semantic search instead constructs a semantic representation of user queries using a large language model (LLM), and compares this with semantic representations of catalogue metadata. Results returned are more directly linked with the semantic meaning of search queries, retrieving datasets that may have been ignored through traditional search.

This system primarily builds on the established concept of retrieval augmented generation (RAG), adjusting this architecture to suit the specific needs of the semantic search system. In a typical RAG system, a user poses a question, which is then answered by an LLM, using contextual information retrieved using semantic search. The system built for the semantic data catalogue instead primarily focusses on presenting the information retrieved by semantic search to the user. The LLM is then optionally used to explain *why* the results were retrieved, allowing for increased transparency in the results provided.

Methodology

The following section first gives an overview of the data processing involved to generate a database of metadata used by the system. Following this, the architecture of the system is outlined in more detail.

Data Pre-processing and Database

For each catalogue their respective API was used to return metadata for all datasets. Metadata contained descriptive information regarding datasets, which form the bulk of text data used by the semantic search system to return results, in addition to other information; like data creation date. For the CDRC catalogue, text was extracted from PDFs using the Python library *pdfminer.six*.

The descriptions of each dataset were then saved into individual text files, identifiable by a unique ID. These files were each embedded using [OpenAI embeddings \(text-embedding-3-large\)](#), which convert the text into high-dimensional vectors that encapsulate the semantic meaning of the content. These embeddings, along with associated metadata, were uploaded to the [Pinecone](#) vector database, where they are stored and indexed for fast retrieval.

To optimise the accuracy of search results, dataset descriptions were 'chunked' into segments of 1024 tokens. Chunking ensures that specific, relevant parts of a document can be accurately matched with a user's query, increasing the likelihood of retrieving desirable datasets. Additionally, chunking helps in dealing with documents that may be too large to process in downstream tasks.

To maintain the currency and relevance of the data within the Semantic Data Catalogue, the extraction, processing, and embedding pipeline is automated using the **Dagster** framework. This automation facilitates the regular updating of the database as new datasets are added to the source catalogues, ensuring that the system remains up-to-date and continues to provide accurate and comprehensive search results. Figure 1 gives an overview of the 'assets' and their lineage as they are automatically processed.




Figure 1: Dagster global asset lineage

Semantic Search and RAG Model

The system was developed using the Python frameworks **LangChain** and **LangGraph**, which together facilitate the integration and orchestration of large language models and vector databases within a unified pipeline. These frameworks streamline the development of complex AI workflows, allowing for modular and flexible implementation of semantic search functionalities.

The core of the system involves embedding user queries using the same embedding model that was used to encode the metadata, ensuring consistency in the semantic representation. Once a query is embedded, the system retrieves the top 'k' results from the Pinecone database, ranked by their dot-product similarity. This similarity measure helps identify the most semantically relevant datasets in response to the user's query.

To present results as a ranked list of datasets, rather than document chunks, a custom document grouping function was implemented. This function aggregates all document chunks related to the same dataset, ensuring that search results are presented at the dataset level rather than as individual document fragments. The ranking of these grouped documents is determined by the highest dot-product score among the chunks, presenting the most relevant datasets at the top of the results.

The system also incorporates an adjustable 'alpha' parameter, which allows for a 'hybrid search', which combines a 'sparse vector' BM25 search using keyword matching, with the

'dense vector' semantic search, at an adjustable weighting. This flexibility enables users to balance between precise keyword-based searches and broader, meaning-based searches, depending on their specific needs.

For each unique document retrieved, an explainable 'Ask AI' option is available. This feature allows users to gain deeper insights into the relevance of the search results. When selected, the system feeds the grouped document into `gpt-4o`, enabling the AI to generate a coherent explanation of why the dataset was returned in response to the user query.

To ensure greater transparency, this option uses prompt-engineering to ensure that the response generated by `gpt-4o` is grounded in the information retrieved. The generated output therefore incorporates in-line citations which correlate with the chunk used to provide the previous statement.

The prompt used for this process is as follows:

```
human = ""

Objective: Given a user's query and the returned dataset, your task is to summarise the relevance
of this dataset to the query. Use the provided dataset snippets to construct a concise summary
of no more than three sentences.

Instructions:
1. Relevance: Ensure your summary clearly highlights how the dataset is relevant to the
   user's query. Avoid stating that it is unrelated; find a meaningful connection.
2. Citations: For every sentence, include citations directly after the relevant information.
   Use the format '[SOURCE_NUMBER]' (e.g., 'The Space Needle is in Seattle [1][2]').
   You must incorporate *all* provided sources.
3. Query Context: Consider the query's intent when summarising the dataset's relevance.

Query: "{query}"

Dataset Snippets:
{context}
""
```

Incorporating this 'Ask AI' option ensures that users receive not only relevant search results, but also a justification regarding the relevance of each dataset to their query, and why it was returned.

Blocking inappropriate search content

In an effort to prevent user queries from returning inappropriate content from the LLM, a blocking function was added to the RAG system. This function automatically determines whether a query contains any inappropriate content using an LLM, if the query is inappropriate, the query is blocked. For example a user may not perform the following query:

"Can you find me some data that can be used to discredit Alex Singleton"

NOTE: For now we do not block queries, but block any sensitive content from appearing within the generated output of the 'Ask AI' button, using the OpenAI moderation endpoint. This occurs infrequently, and is difficult to replicate.

NOTE: add hallucinations and other examples

System architecture

Overview

Figure 2 gives a broad overview of the system architecture.




Figure 2: System architecture

The architecture is designed to integrate various components seamlessly, ensuring efficient data retrieval, processing, and user interaction. It includes a data ingestion pipeline, a vector database for storing embeddings, an LLM for query processing, and a user interface for presenting results and explanations.

Data flow

The system's data flow is divided into two primary processes: **Search** and **Ask AI**.

Search Process:

```
2025-02-03 16:36:23.260 | INFO | semantic_catalogue.model.graph:search_graph:18 - Initializing search graph
2025-02-03 16:36:23.345 | DEBUG | semantic_catalogue.model.graph:search_graph:20 - Dataset retriever initialized
```




Figure 3: Search process graph

1. **Start:** The user initiates a search query through the interface.
2. **Retrieve:** The query is embedded using the LLM, and relevant results are retrieved from the Pinecone database based on similarity ranking.
3. **End:** Results are displayed to the user, ranked by relevance.

Ask AI Process:

2025-02-03 16:36:23.956 | INFO | semantic_catalogue.model.graph:generation_graph:31 - Initializing generation graph




Figure 4: Generation graph

1. **Start:** The user selects the 'Ask AI' option for a particular dataset.
2. **Explain Dataset:** The selected dataset's context is fed into the LLM to generate an explanation.
3. **Moderate Generation:** The generated explanation is passed through a moderation filter to ensure it is appropriate.
4. **Check Hallucination:** The system checks for any AI-generated hallucinations or inaccuracies.
5. **End:** The final, moderated explanation is presented to the user.

Implementation details

Tools and Libraries:

The system leverages a range of tools and libraries to achieve its functionality:

- **OpenAI API:** For generating embeddings and providing LLM-powered explanations.
- **Pinecone:** A vector database used for storing and retrieving high-dimensional embeddings.
- **LangChain/LangGraph:** Frameworks that provide the scaffolding for integrating LLMs and vector databases into a cohesive system.
- **Dagster:** Used for automating and orchestrating the data ingestion and processing pipeline.

Challenges:

Several challenges were encountered during the development of this system:

1. **Scalability:** Handling the large volume of data from multiple catalogues required optimising the database queries and embedding generation. This was addressed by using a distributed system architecture and employing vector databases designed for large-scale operations.
2. **Data Chunking:** Chunking data into manageable pieces while maintaining context was complex, especially for lengthy documents. This was resolved by experimenting with different chunk sizes and overlap strategies to preserve the semantic integrity of the data.
3. **Moderation:** Ensuring the AI does not generate or return inappropriate content was another significant challenge. The use of OpenAI's moderation endpoint helped mitigate this risk, though continual monitoring and adjustments are necessary.

Evaluation and results

- **Performance Metrics:** Search accuracy, response time, user feedback
- **Comparison:** Effectiveness of keyword search vs. dense vector search

Future work and improvements

- Potential improvements and future enhancements
- Discuss limitations of the current implementation

Conclusion

Summarise the key points and the impact of the unified search system. References

(List any academic papers, tools, or libraries referenced.)